

# Lecture 20 – NVHPC Hands-on

Prof. Brendan Kochunas

NERS/ENGR 570 - Methods and Practice of Scientific Computing (F22)



COLLEGE OF ENGINEERING  
NUCLEAR ENGINEERING & RADIOLOGICAL SCIENCES  
UNIVERSITY OF MICHIGAN

# Outline

- Announcements
- (Labs 9 and 10) Simulated Annealing Problem with OpenMP
- (HW3) SpMV with CUDA

# Learning Objectives: By the end of Today's Lecture you should be able to

- (*Skill*) write a better OpenMP implementation of Simulated Annealing
- (*Skill*) use the OpenMP TARGET directive for GPU offload
- (*Skill*) write simple CUDA kernels (maybe)



# Simulated Annealing

# Instructions

- Pair/Group up (so we don't have a bunch simultaneous jobs on GL)
  - Login to (Log into? Log in to?) Great Lakes
  - Copy down laplace\_serial.c into new directory for today's lecture
- Establish host Scaling
  - Login to (Log into? Log in to?)
  - `$ srun -A ners570f22 class --pty --partition=standard --time=00:60:00 --nodes=1 --ntasks-per-node=1 --cpus-per-task=36 --mem=100gb /bin/bash`
- Work on GPU implementations
  - Start a new interactive job on GL
  - `$ srun -A ners570f22 class --pty --partition=gpu --time=00:60:00 --gpus=1 /bin/bash`

# My Host Speedup

Number of Threads	gcc -Ofast -march=native -fopenmp (Time s)	nvc -fast -mp (Time s)	GCC Speedup	NVC Speedup
1	5.101	5.556	1.0	1.0
2	2.325	2.591	2.2	2.1
4	1.181	1.277	4.3	4.4
8	0.710	0.776	7.2	7.2
16	0.382	0.370	13.4	15.0
24	0.155	0.192	32.9	28.9
32	0.129	0.163	39.6	34.0
36	0.125	0.159	40.9	35.1

# My OpenMP Solution

```
59 // do until error is minimal or until max steps
60 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
61
62     dt = 0.0; // reset largest temperature change
63 #pragma omp parallel default(shared) private(i,j)
64 {
65
66     // main calculation: average my four neighbors, compute local max dt
67     #pragma omp for
68     for(i = 1; i <= ROWS; i++) {
69         for(j = 1; j <= COLUMNS; j++) {
70             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
71                                         Temperature_last[i][j+1] + Temperature_last[i][j-1]);
72         }
73     }
74
75     // copy grid to old grid for next iteration
76     #pragma omp for reduction(max:dt)
77     for(i = 1; i <= ROWS; i++){
78         for(j = 1; j <= COLUMNS; j++){
79             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
80             Temperature_last[i][j] = Temperature[i][j];
81         }
82     }
83 }
84
85 // periodically print test values
86 if((iteration % 100) == 0) {
87     track_progress(iteration);
88 }
89
90 iteration++;
91 }
```

# Now we do GPUs

```
$ srun -A ners570f22_class --pty --partition=gpu --  
time=00:60:00 --gpus=1 /bin/bash
```



# First GPU Implementation (Threaded)

```
59 // do until error is minimal or until max steps
60 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
61     dt = 0.0; // reset largest temperature change
62
63     #pragma omp target map(Temperature_last) map(Temperature) map(dt)
64     {
65         // main calculation: average my four neighbors, compute local max dt
66         #pragma omp parallel for
67         for(i = 1; i <= ROWS; i++) {
68             for(j = 1; j <= COLUMNS; j++) {
69                 Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
70                                             Temperature_last[i][j+1] + Temperature_last[i][j-1]);
71             }
72         }
73
74         // copy grid to old grid for next iteration
75         #pragma omp parallel for reduction(max:dt)
76         for(i = 1; i <= ROWS; i++){
77             for(j = 1; j <= COLUMNS; j++){
78                 dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
79                 Temperature_last[i][j] = Temperature[i][j];
80             }
81         }
82     }
83 }
84 // periodically print test values
85 if((iteration % 100) == 0) {
86     track_progress(iteration);
87 }
88
89 iteration++;
90 }
91 }
```

# GPU With Teams

```
59 #pragma omp target data map(Temperature) map(Temperature_last)
60 // do until error is minimal or until max steps
61 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
62
63     // main calculation: average my four neighbors, compute local max dt
64 #pragma omp target teams distribute parallel for
65     for( i = 1; i <= ROWS; i++) {
66         for( j = 1; j <= COLUMNS; j++) {
67             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
68                                         Temperature_last[i][j+1] + Temperature_last[i][j-1]);
69         }
70     }
71
72     dt = 0.0; // reset largest temperature change
73
74     // copy grid to old grid for next iteration
75 #pragma omp target teams distribute parallel for reduction(max:dt)
76     for( i = 1; i <= ROWS; i++){
77         for( j = 1; j <= COLUMNS; j++){
78             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
79             Temperature_last[i][j] = Temperature[i][j];
80         }
81     }
82
83     // periodically print test values
84     if((iteration % 100) == 0) {
85         track_progress(iteration);
86     }
87
88     iteration++;
89 }
```

Explicitly maps arrays  
For entire while loop

Spawns thread teams  
And distributes iterations  
to those teams


Worksharing within teams

# The Teams Directive

## TEAMS Directive

To better utilize the GPU resources, use many thread teams via the TEAMS directive.

- Spawns 1 or more thread teams with the same number of threads
- Execution continues on the master threads of each team (redundantly)
- No synchronization between teams

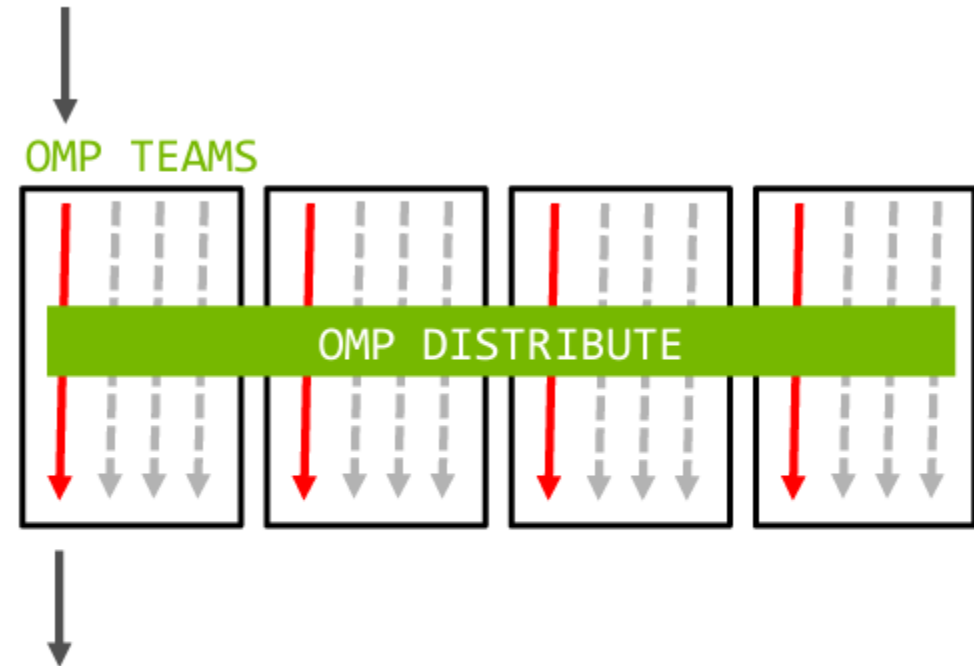


# The Distribute Directive

## DISTRIBUTE Directive

Distributes the iterations of the next loop to the master threads of the teams.

- Iterations are distributed statically.
- There's no guarantees about the order teams will execute.
- No guarantee that all teams will execute simultaneously
- Does not generate parallelism/worksharing within the thread teams.



# GPU Teams with Splitting loops

```
59 #pragma omp target data map(Temperature) map(Temperature_last)
60 // do until error is minimal or until max steps
61 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
62     // main calculation: average my four neighbors, compute local max dt
63     #pragma omp target teams distribute
64     for(i = 1; i <= ROWS; i++) {
65         #pragma omp parallel for
66         for(j = 1; j <= COLUMNS; j++) {
67             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
68             Temperature_last[i][j+1] + Temperature_last[i][j-1]);
69         }
70     }
71     dt = 0.0; // reset largest temperature change
72     // copy grid to old grid for next iteration
73     #pragma omp target teams distribute reduction(max:dt)
74     for(i = 1; i <= ROWS; i++){
75         #pragma omp parallel for reduction(max:dt)
76         for(j = 1; j <= COLUMNS; j++){
77             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
78             Temperature_last[i][j] = Temperature[i][j];
79         }
80     }
81     // periodically print test values
82     if((iteration % 100) == 0) {
83         track_progress(iteration);
84     }
85     iteration++;
86 }
```

Distribute “i” loops over teams

Workshare “j” loops over threads

# GPU Teams with loop collapsing

```
59 #pragma omp target data map(Temperature) map(Temperature_last)
60 // do until error is minimal or until max steps
61 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
62
63     // main calculation: average my four neighbors compute local max dt
64     #pragma omp target teams distribute parallel for collapse(2)
65     for(i = 1; i <= ROWS; i++) {
66         for(j = 1; j <= COLUMNS; j++) {
67             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
68                                         Temperature_last[i][j+1] + Temperature_last[i][j-1]);
69         }
70     }
71
72     dt = 0.0; // reset largest temperature change
73
74     // copy grid to old grid for next iteration
75     #pragma omp target teams distribute parallel for reduction(max:dt) collapse(2)
76     for(i = 1; i <= ROWS; i++){
77         for(j = 1; j <= COLUMNS; j++){
78             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
79             Temperature_last[i][j] = Temperature[i][j];
80         }
81     }
82
83     // periodically print test values
84     if((iteration % 100) == 0) {
85         track_progress(iteration);
86     }
87
88     iteration++;
89 }
```

Transforms 2-loops  
Into 1-loop for  
Increased parallelism

# GPU Teams with loop Splitting and Scheduling

```
59 #pragma omp target data map(Temperature) map(Temperature_last)
60 // do until error is minimal or until max steps
61 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
62
63     // main calculation: average my four neighbors, compute local max dt
64 #pragma omp target teams distribute
65     for(i = 1; i <= ROWS; i++) {
66 #pragma omp parallel for schedule(static,1)
67         for(j = 1; j <= COLUMNS; j++) {
68             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
69                                         Temperature_last[i][j+1] + Temperature_last[i][j-1]);
70         }
71     }
72
73     dt = 0.0; // reset largest temperature change
74
75     // copy grid to old grid for next iteration
76 #pragma omp target teams distribute reduction(max:dt)
77     for(i = 1; i <= ROWS; i++){
78 #pragma omp parallel for reduction(max:dt) schedule(static,1)
79         for(j = 1; j <= COLUMNS; j++){
80             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
81             Temperature_last[i][j] = Temperature[i][j];
82         }
83     }
84
85     // periodically print test values
86     if((iteration % 100) == 0) {
87         track_progress(iteration);
88     }
89
90     iteration++;
91 }
```

Align loop iterations  
with threads

# GPU Teams with loop collapsing and Scheduling


```
59 #pragma omp target data map(tofrom:Temperature) map(Temperature_last)
60 // do until error is minimal or until max steps
61 while ( dt > MAX_TEMP_ERROR && iteration <= max_iterations ) {
62
63     // main calculation: average my four neighbors, compute local max dt
64 #pragma omp target teams distribute parallel for collapse(2) schedule(static,1)
65     for(i = 1; i <= ROWS; i++) {
66         for(j = 1; j <= COLUMNS; j++) {
67             Temperature[i][j] = 0.25 * (Temperature_last[i+1][j] + Temperature_last[i-1][j] +
68                                         Temperature_last[i][j+1] + Temperature_last[i][j-1]);
69         }
70     }
71
72     dt = 0.0; // reset largest temperature change
73
74     // copy grid to old grid for next iteration
75 #pragma omp target teams distribute parallel for reduction(max:dt) collapse(2) schedule(static,1)
76     for(i = 1; i <= ROWS; i++){
77         for(j = 1; j <= COLUMNS; j++){
78             dt = fmax( fabs(Temperature[i][j]-Temperature_last[i][j]), dt);
79             Temperature_last[i][j] = Temperature[i][j];
80         }
81     }
82
83     // periodically print test values
84     if((iteration % 100) == 0) {
85         track_progress(iteration);
86     }
87
88     iteration++;
89 }
90
```

Align loop iterations  
with threads




# Summary of Speedup

1000x1000 Speedup (Higher is Better)



# What if we have a Larger Problem?

4096x4096 Speedup (Higher is Better)





# SpMV ELLPACK Format