

Lecture 24 – Debugging and Profiling Tools

Prof. Brendan Kochunas

NERS/ENGR 570 - Methods and Practice of Scientific Computing (F22)



COLLEGE OF ENGINEERING

NUCLEAR ENGINEERING & RADIOLOGICAL SCIENCES

UNIVERSITY OF MICHIGAN

Outline

- Overview of Debugging
- Debugging Tools
 - Help from the Compiler
 - From the Command Line (GDB)
 - Memory Debugging (Valgrind)
 - Fancy Debugging Tools (ARM DDT)
- Overview of Profiling
 - Types of Performance Data
 - Collecting Performance Data
 - Analyzing Performance Data
- Performance Analysis Tools

Learning Objectives: By the end of Today's Lecture you should be able to

- (*Knowledge*) Classes of bugs and observed behaviors
- (*Skill*) Debugging at the most basic level
- (*Knowledge*) What types of tools are available for debugging
- (*Skill*) debugging with gdb, valgrind, and ddt


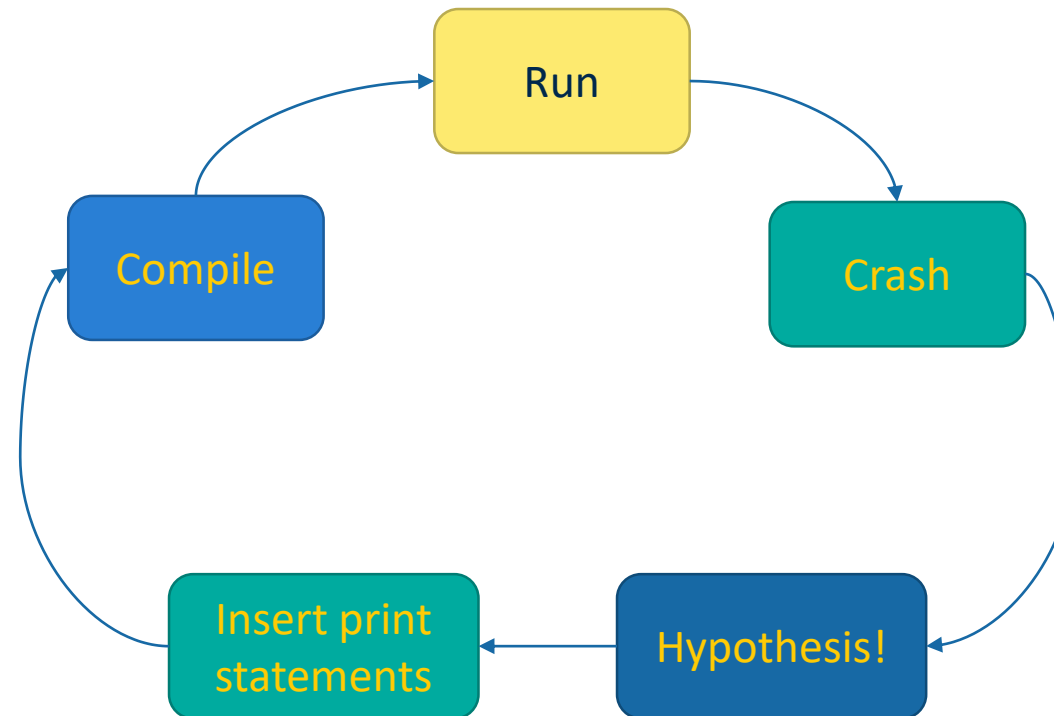



Overview of Debugging

What is debugging?

- The art of transforming a broken program to a working one
- Debugging requires thought – and discipline:

Debugging in Practice...



Bang Head Here

A Taxonomy of Bugs

Name	Description
<i>Bohrbug</i>	Steady, dependable bug
<i>Heisenbug</i>	Vanishes when you try to debug (observe)
<i>Mandelbug</i>	Complexity and obscurity of the cause is so great that it appears chaotic
<i>Schroedinbug</i>	First occurs after someone reads the source file and deduces that it never worked, after which the program ceases to work

From presentation: David Lecomber, “Debugging and Profiling your HPC Applications”

Common Approaches to Debugging



From presentation: David Lecomber, “Debugging and Profiling your HPC Applications”



Compiler Supported Debugging

What compilers are good for in debugging

- Catching typos
 - Gives a compile time error
- Enforcing variable definitions/type consistency
- Adding instrumentation
 - Checking array bounds
- Providing callstack information
 - backtrace
- Can add warnings, and make these errors
- Integrating source code information for other debugging tools

Summary of Options

GCC compiler option	Meaning
<code>-g</code>	Produce debugging information in the operating system's native format.
<code>-fbounds-check</code>	Generate code to check that indices used to access arrays are within the declared range during run time.
<code>-fstack-check</code>	Generate code to verify that you do not go beyond the boundary of the stack.
<code>-fbacktrace</code>	Prints the call stack when there are run time errors
<code>-fcheck-pointer-bounds</code>	Each memory reference is instrumented with checks of the pointer used for memory access against bounds associated with that pointer.
<code>-fsanitize=<opt></code>	Enable AddressSanitizer, a fast memory error detector.
<code>-Wall</code>	Enable all warnings
<code>-Wpedantic</code>	Issue all the warnings demanded by strict ISO C and ISO C++
<code>-Werror</code>	Treat warnings as compiler errors



GNU Debugger (GDB)

Debugging with GDB

- Invoked with `$ gdb <exec> [<exec_args>...]`
 - Generates a new command line interface
- A few basic commands...

Command	Description
help	Display help
break	Add a breakpoint
delete	Remove breakpoint(s)
continue	Run program if stopped (till next breakpoint)
step	Executes one line of source
next	Like step, but does not enter functions
list <linenumber>	Print source near line number
print <expression>	Print the result of an expression



Memory Debugging

Memory Debugging with Valgrind



- Essentially a “virtual machine” that simulates all memory traffic and catches errors.
 - Actually a collection of tools, but we’ll focus on memcheck
- Primarily used for serial code
 - Difficult to scale to large applications
- <http://www.valgrind.org>

Basic Usage

1. Compile your software with debug symbols (e.g. `-g`)
2. Rerun your program with the command:

```
$ valgrind <myexe> [<myexe_arguments>...]
```

Types errors detected by Valgrind

- Memory leaks
 - Memory is allocated but not deallocated
 - Primarily happens with pointers
 - Easier to do in C/C++
 - Serious problem if performed in a loop
- Illegal read/writes
 - When program accesses illegal memory addresses (e.g. beyond the bounds of an array).
- Use of uninitialized values
 - Serious if this happens to be in a branching statement



Fancy Debuggers


List of “Fancy” Debuggers

- Various GDB “front ends”
 - atom, vscode, gdbgui, Eclipse, Kdevelop
- DDT – Best option for leadership class computers
 - This is the best, but costs \$\$\$\$. Fortunately, most (good) computing centers have this installed
- Totalview – Used to be best tool, but likely surpassed by DDT
 - Costs \$\$\$ some computing centers have it
- STAT – Stack Trace Analysis Tool
 - Free! <https://github.com/LLNL/STAT>



Overview of Profiling

Why do Profiling?



<https://xkcd.com/1691/>



THE REASON I AM SO INEFFICIENT

<https://xkcd.com/1445>



Types of Performance Data



Metrics (Schmetrics)

Primitive

Derived



Profiling vs Tracing

Profile

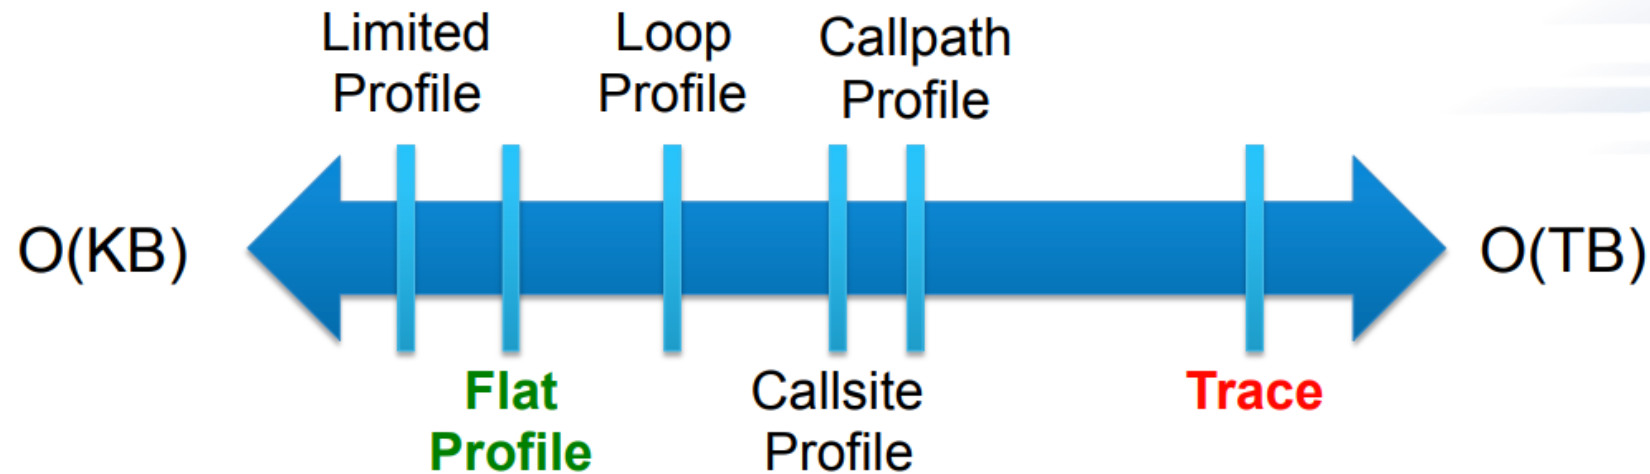
Trace

Inclusive vs Exclusive

ALGORITHM 6.9 GMRES

1. Compute $r_0 = b - Ax_0$, $\beta := \|r_0\|_2$, and $v_1 := r_0/\beta$
2. For $j = 1, 2, \dots, m$ Do:
3. Compute $w_j := Av_j$
4. For $i = 1, \dots, j$ Do:
5. $h_{ij} := (w_j, v_i)$
6. $w_j := w_j - h_{ij}v_i$
7. EndDo
8. $h_{j+1,j} = \|w_j\|_2$. If $h_{j+1,j} = 0$ set $m := j$ and go to 11
9. $v_{j+1} = w_j/h_{j+1,j}$
10. EndDo
11. Define the $(m+1) \times m$ Hessenberg matrix $\bar{H}_m = \{h_{ij}\}_{1 \leq i \leq m+1, 1 \leq j \leq m}$.
12. Compute y_m the minimizer of $\|\beta e_1 - \bar{H}_m y\|_2$ and $x_m = x_0 + V_m y_m$.

How much data do you want?





Collecting Performance Data

Instrumenting Code

ALGORITHM 6.9 GMRES

1. Compute $r_0 = b - Ax_0$, $\beta := \|r_0\|_2$, and $v_1 := r_0/\beta$
2. For $j = 1, 2, \dots, m$ Do:
3. Compute $w_j := Av_j$
4. For $i = 1, \dots, j$ Do:
5. $h_{ij} := (w_j, v_i)$
6. $w_j := w_j - h_{ij}v_i$
7. EndDo
8. $h_{j+1,j} = \|w_j\|_2$. If $h_{j+1,j} = 0$ set $m := j$ and go to 11
9. $v_{j+1} = w_j/h_{j+1,j}$
10. EndDo
11. Define the $(m+1) \times m$ Hessenberg matrix $\bar{H}_m = \{h_{ij}\}_{1 \leq i \leq m+1, 1 \leq j \leq m}$.
12. Compute y_m the minimizer of $\|\beta e_1 - \bar{H}_m y\|_2$ and $x_m = x_0 + V_m y_m$.

Overhead from Instrumentation

ALGORITHM 6.9 GMRES

1. Compute $r_0 = b - Ax_0$, $\beta := \|r_0\|_2$, and $v_1 := r_0/\beta$
2. For $j = 1, 2, \dots, m$ Do:
3. Compute $w_j := Av_j$
4. For $i = 1, \dots, j$ Do:
5. $h_{ij} := (w_j, v_i)$
6. $w_j := w_j - h_{ij}v_i$
7. EndDo
8. $h_{j+1,j} = \|w_j\|_2$. If $h_{j+1,j} = 0$ set $m := j$ and go to 11
9. $v_{j+1} = w_j/h_{j+1,j}$
10. EndDo
11. Define the $(m+1) \times m$ Hessenberg matrix $\bar{H}_m = \{h_{ij}\}_{1 \leq i \leq m+1, 1 \leq j \leq m}$.
12. Compute y_m the minimizer of $\|\beta e_1 - \bar{H}_m y\|_2$ and $x_m = x_0 + V_m y_m$.

Automatic Instrumentation and Sampling



Scaling Studies

- Strong Scaling: fixes problem size and increases number of processors.
 - Provides insight into how finely grained an algorithm can be parallelized and how much parallel overhead there is relative to useful computation
- Weak Scaling: fixes problem size *per process* and increases number of processors.
 - Provides insight into whether the parallel overhead varies faster or slower than the amount of work as the problem size is increased.

- Speedup and Efficiency:
$$S(P_{size}, N_p) \equiv \frac{T(P_{size}, 1)}{T(P_{size}, N_p)}$$

$$E_{strong}(P_{size}, N_p) \equiv \frac{T(P_{size}, 1)}{N_p \times T(P_{size}, N_p)}$$

- Good efficiency does not necessarily mean you have fast code
 - It could mean you have terrible serial performance

$$E_{weak}(P_{size}, N_p) \equiv \frac{T(P_{size}, 1)}{T(P_{size} \times N_p, N_p)}$$



Performance Analysis



How to interpret your data?

Low Computational Intensity

Simple Performance Model

$$T = Ft_F \left(1 + \frac{t_M}{t_F} \frac{L}{F} \right)$$

F = # of FLOPs

L = # of loads and stores


t_F = time for flop

t_M = time for memory load/store

T = execution time

$$\frac{L}{F} = \frac{1}{q}$$


Roofline Model



Load Balance

Problem

- Poor strong scaling
- Poor parallel efficiency
- Increase cores by factor of 2x, do not observe 2x speedup.



Solution

- Determine what the load balance/imbalance is
 - Need to assign a value of “work” to each subdomain.
 - What is the maximum to minimum workload for all domains.
- Change partitioning to improve load balancing
- Change parallel algorithm

Parallel Execution Time Models

Moving from serial to parallel

- Serial Latency based model

$$T_{serial} = Ft_F + \alpha_1 L + \sum_{j=1}^{K-1} (\alpha_{j+1} - \alpha_j) M_j + (\alpha_{mem} - \alpha_K) M_K$$

- Parallel Model

$$T_{parallel}(N_p) = \frac{T_{serial}}{N_p} + T_{overhead}(N_p)$$

- Difficult to develop exact expressions,
 - Alternatively measure realistic average values based on microbenchmarks.

Canonical Execution Time Models

- Distributed Memory Computing
 - Point-to-Point Communication Time

$$T_{comm} = \alpha_{network} + \beta_{network} N$$

Latency \uparrow $\frac{1}{\text{Bandwidth}}$ \uparrow Amount of data

- Collective operations have their own (depends on algorithm implemented in library)

$$T_{All_reduce, small} = \lceil \log p \rceil (\alpha_{network} + \beta_{network} \times N + \gamma \times N)$$

$$T_{All_reduce, large} = 2 \log p \alpha_{network} + \frac{p-1}{p} (2 \beta_{network} \times N + \gamma \times N)$$

Time to perform reduce operation
(e.g. sum, max, multiply, etc.)

Fundamentals of getting good parallel performance

- Maximize amount of work that can be parallelized.
- Minimize overhead.
- Usually this means
 - Balance work loads among processors
 - Avoid synchronization
 - Especially for shared memory
 - use non-blocking communication
 - Primarily in distributed memory models
- Make sure the serial code is optimized.

Assumes perfect load balance

$$T_{parallel}(N_p) = \underbrace{T_{non-parallel}}_{\text{Minimize}} + \overbrace{\frac{T_{serial}}{N_p}}^{\text{Assumes perfect load balance}} + \underbrace{T_{overhead}(N_p)}_{\text{Minimize}}$$



Performance Analysis Tools



Performance Utilities and Tools